

SCIT

School of Computing & Information Technology

CSCI376 – Multicore and GPU Programming SIM Session 2 2020

Assignment 2

Task 1 – OpenCL Vector Datatype (5 marks)

Write a program that does the following:

- In the host, create two arrays as follows:
 - Array 1: An 8 element array of ints with random values between 10 and 20.
 - Array 2: A 16 element array of ints. Initialise the first half of the array with values from 2 to 9 and the second half with values from -9 to -2.

(1 mark)
- Write a kernel that
 - Accepts array 1 as an array of int4s, array 2 and an output array
 - Reads the contents from array 1 and 2 into local memory
 - Copy the contents of array 1 into an int8 vector called *v*
 - Copy the contents (using **vloadn**) of array 2 into two int8 vectors called *v1* and *v2*
 - Creates an int8 vector in private memory called *results*. The contents of this vector should be filled as follows:
 - Check whether **any** of the elements in *v* are greater than 15
 - If there are, then for elements that are greater than 15, copy the corresponding elements from *v1* into *results*; for elements less than or equal to 15, copy the elements from *v2* into *results*. (Use **select**).
 - If not, fill the first 4 elements of *results* with the contents from the first 4 elements of *v1*; and fill the next 4 elements of *results* with contents from the first 4 elements of *v2*.
 - Stores the contents of *v*, *v1*, *v2* and *results* in the output array (using **vstoren**)

Note that the host will only have to enqueue 1 work item.

(3 marks)

- In the host, check that the results are correct and display the contents of the output array.

(1 mark)

Task 2 – Shift Cipher (10 marks)

A shift cipher (a.k.a. Caesar's cipher) is a simple substitution cipher in which each letter in the plaintext is replaced with another letter that is located a certain number, n , positions away in the alphabet. The value of n can be positive or negative. For positive values, replace letters with letters located n places on its right (i.e. 'shifted' by n positions to the right). For negative values, replace letters with letters located n places on its left. If it reaches the end/start of the alphabets, wrap around to the start/end.

For example: If $n = -3$, each letter in the plaintext is replaced with a letter 3 positions before that letter in the alphabet list.

Plaintext: **The quick brown fox jumps over the lazy dog.**

Ciphertext: **QEB NRFZH YOLTK CLU GRJMP LSBO QEB IXWV ALD.**

Note that in the example above, $c \rightarrow Z$, since 3 positions before 'c' wraps around to the end of the alphabet list and continues from 'Z'. Similarly, $a \rightarrow X$ and $b \rightarrow Y$.

Leave anything that is not an alphabet as is (i.e. punctuations and spaces).¹

Decrypting the ciphertext is simply a matter of reversing the shift.

Task 2a

Write a normal C/C++ program (not using OpenCL) that reads the contents from a text file called "plaintext.txt" (a test file has been provided). The program should prompt the user to input a valid n value, then encrypt the plaintext using the shift cipher method described above, and output the ciphertext into an output text file called "ciphertext.txt". To ensure that the encryption was performed correctly, your program must also decrypt the ciphertext into a file called "decrypted.txt" to check whether it matches the original plaintext (albeit in upper case).

(3 marks)

Task 2b

Write an OpenCL program to perform the same functionality as in Task 2a, but in parallel. Note that it is more efficient to use OpenCL vector datatypes for processing in the kernel, as less work-items will be required.

(3 marks)

Task 2c

Write an OpenCL program to perform parallel encryption, and decryption, by substituting characters based on the following lookup table:

a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
G	X	S	Q	F	A	R	O	W	B	L	M	T	H	C	V	P	N	Z	U	I	E	Y	D	K	J

Based on the table above, for encryption the letter a (or A) will be replaced by G, b (or B) will be replaced by X, c (or C) will be replaced by S, etc.

(4 marks)

¹ Note that to avoid leaking information (e.g., word length), by convention the ciphertext is usually converted to upper case letters that are organised in groups of five-letter blocks and anything that is not a letter is removed. However, for this assignment, DO NOT remove anything that is not a letter and DO NOT organise in groups of five-letters.

Task 3 – Parallel Image Processing (10 marks)

For this section, a test image, “peppers.bmp”, has been provided.

Task 3a (Image Luminance)

Write a parallel program to convert the RGB values (i.e. Red, Green and Blue colour channels) in an image to luminance values (this approach is used to convert a colour image into a greyscale image).

For each pixel, calculate:

$$\text{Luminance} = 0.299 * R + 0.587 * G + 0.114 * B$$

Save the luminance image into a 24-bit BMP file. To do this, set the RGB values of each pixel to the luminance value.

Note that if you use the example code from the tutorial on image processing, the R, G, and B values range from 0 to 255 on the host (unsigned char), and 0.0 to 1.0 (float) on the device.

(1 mark)

Task 3b (Gaussian Blurring)

Gaussian blurring is a commonly used technique to image processing and graphics to create a smooth blurring effect using a Gaussian function. The weights of the filter depend on the size of the Gaussian filter window. The following are example weights for a 7x7 windows:


0.000036	0.000363	0.001446	0.002291	0.001446	0.000363	0.000036
0.000363	0.003676	0.014662	0.023226	0.014662	0.003676	0.000363
0.001446	0.014662	0.058488	0.092651	0.058488	0.014662	0.001446
0.002291	0.023226	0.092651	0.146768	0.092651	0.023226	0.002291
0.001446	0.014662	0.058488	0.092651	0.058488	0.014662	0.001446
0.000363	0.003676	0.014662	0.023226	0.014662	0.003676	0.000363
0.000036	0.000363	0.001446	0.002291	0.001446	0.000363	0.000036

- i. Write an OpenCL program that accepts a colour image and outputs a filtered image using Gaussian blurring based on the 7x7 window weights provided above.

(1 mark)

- ii. Instead of using the 7x7 window (the naïve approach), an alternate approach is to run the filter in 2 passes. The first pass will perform blurring in the horizontal direction; the result will then undergo a second pass to blur it in the vertical direction (enqueue the kernel twice to perform blurring in each direction). The result will be similar to the single window approach, but the amount of computation will be different.

For example, using a 7x7 window approach, each pixel will have to perform a weighted sum on 49 pixels. In the 2-pass approach, each pixel will have to perform a weighted sum on 7 pixels in each pass, processing a total of 14 pixels. This is illustrated below:



Your task is to implement the parallel 2-pass approach. For this, use the following weights for the horizontal pass as well as the vertical pass:

0.00598 0.060626 0.241843 0.383103 0.241843 0.060626 0.00598

(3 marks)

Task 3c (Bloom effect)

Bloom effects are commonly used in graphics, movies, video games, etc. This part combines the work from Tasks 3a and 3b. The basic steps to create an image with a bloom effect are illustrated as follows:




Figure 1: Bloom effect steps.

1. The image in Fig. 1(a) shows the original image
2. The image in Fig. 1(b) shows an image where the glowing pixels are kept, while the rest are set to black². For this assignment, allow the user to input a valid threshold luminance value. Pixels above the threshold luminance value are kept, while pixels below this luminance value are set to black. This step is related to Task 3a.
3. The image in Fig. 1(b) undergoes a horizontal blur pass, then a vertical blur pass to obtain the image depicted in Fig. 1(c). This step is related to Task 3b.

² Note that Figure 1(b) shows a down-sampled image (i.e. the image has been shrunk). It is more efficient to process a down-sampled image during the blurring step because there are fewer pixels to process. This also works well for blurring since referencing the pixels from the smaller image in the final step will also cause blurring (blurring caused by effectively up-sampling back to the original image size. In fact, some approaches simply use down-sampling for blurring). To make things easier, for this assignment you do not have to perform down-sampling/up-sampling.

4. Finally, the pixel values in the images shown in Fig. 1(a) and Fig. 1(c) are added together to form the final image shown in Fig. 1(d). Note that the values above the maximum colour value should be clamped to the maximum value.

Write a parallel program in OpenCL to perform the Bloom effect on an input image. For the threshold value (in step 2), allow the user to enter a valid threshold value.

Your program should output the following images:

- an image after step 2 (i.e. image showing the glowing pixels)
- an image after the horizontal blur pass
- an image after the vertical blur pass
- the final image with the bloom effect

(5 marks)

For **ALL** tasks, include **screenshots** with your submission. The screenshots are to demonstrate that the programs work on your computer.

For Task 3, include examples of output images obtained from your program.

Instructions and Assessment

Submit your tasks as one zip file with three folders, named Task1, Task2 and Task3, and include all the required files (e.g., .cpp, .h, .cl) in your submission.

The assignment must be your own work. If asked, you must be able to explain what you did and how you did it. Marks will be deducted if you cannot correctly explain your code. The marking allocations shown above are merely a guide. Marks will be awarded based on the overall quality of your work. Marks may be deducted for other reasons, e.g., if your code is too messy or inefficient, is not well commented, if you cannot correctly explain your code, etc. For code that does not compile, does not work or for programs that crash, the most you can get is half the assessment marks or less.

References

The images were sourced from

- http://http.developer.nvidia.com/GPUGems/gpugems_ch21.html